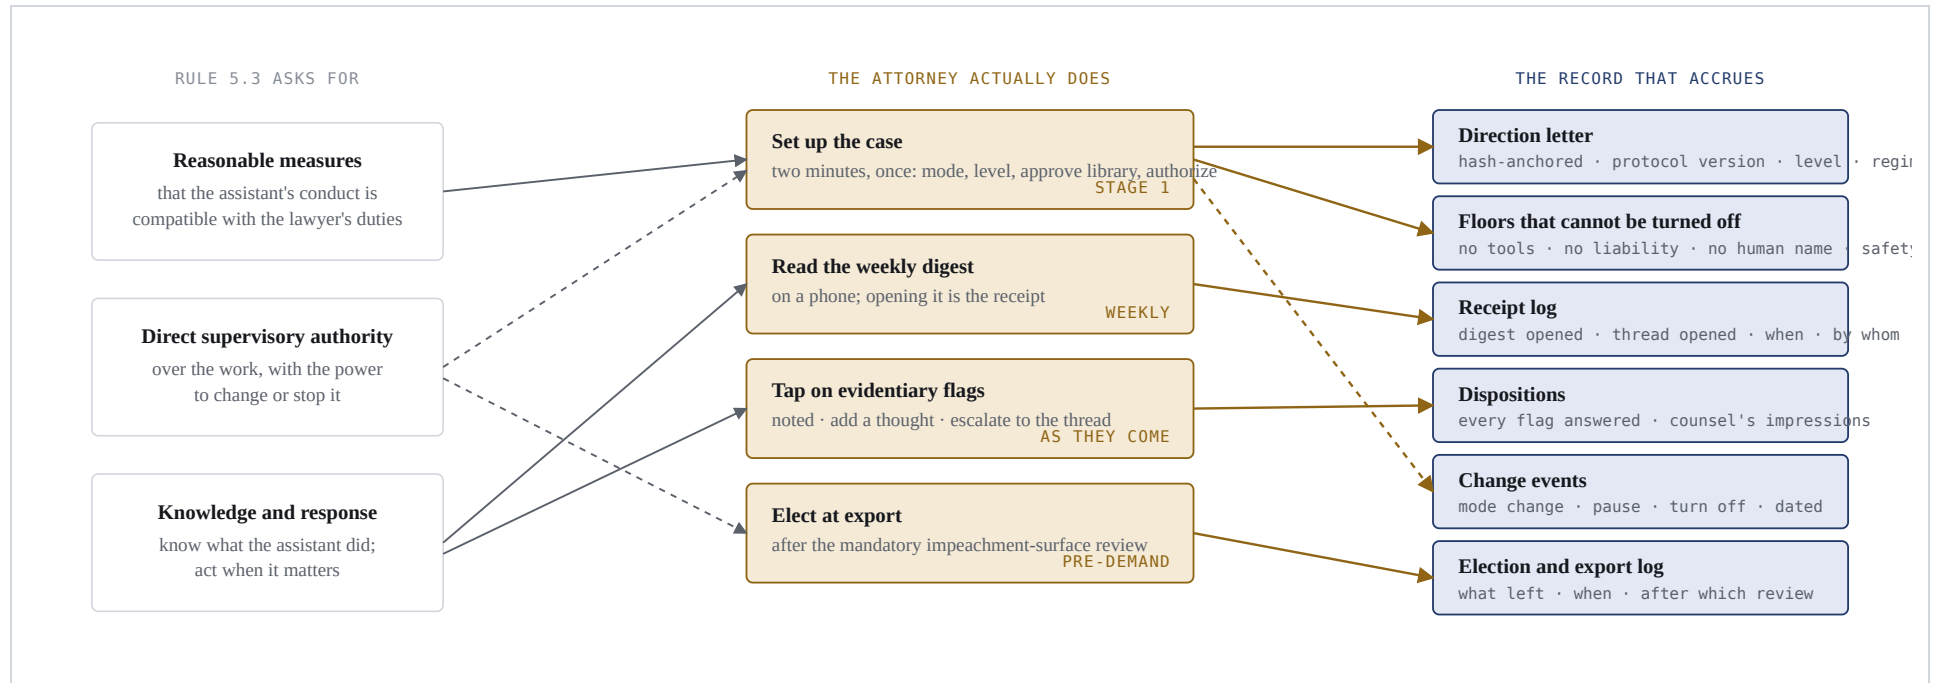


Nobody performs a task called supervision.

Rule 5.3 asks a lawyer for three things about a nonlawyer assistant: reasonable measures, direct supervisory authority, and knowledge of and response to its conduct. PIRRA produces the record of all three **underneath four things the attorney was going to do anyway**. If a supervisory act takes more than one tap, the product is designed wrong.

FIGURE 1

Three duties, four ordinary acts, one record



Setting up the case is the direction and installs the measures. Reading the digest is the knowledge. Tapping a flag is the response, and each tap is also an entry carrying counsel's impressions. Electing at export is the authority exercised. The record on the right exports in one click whenever anyone asks how the assistant was supervised on this case.

FIGURE 2

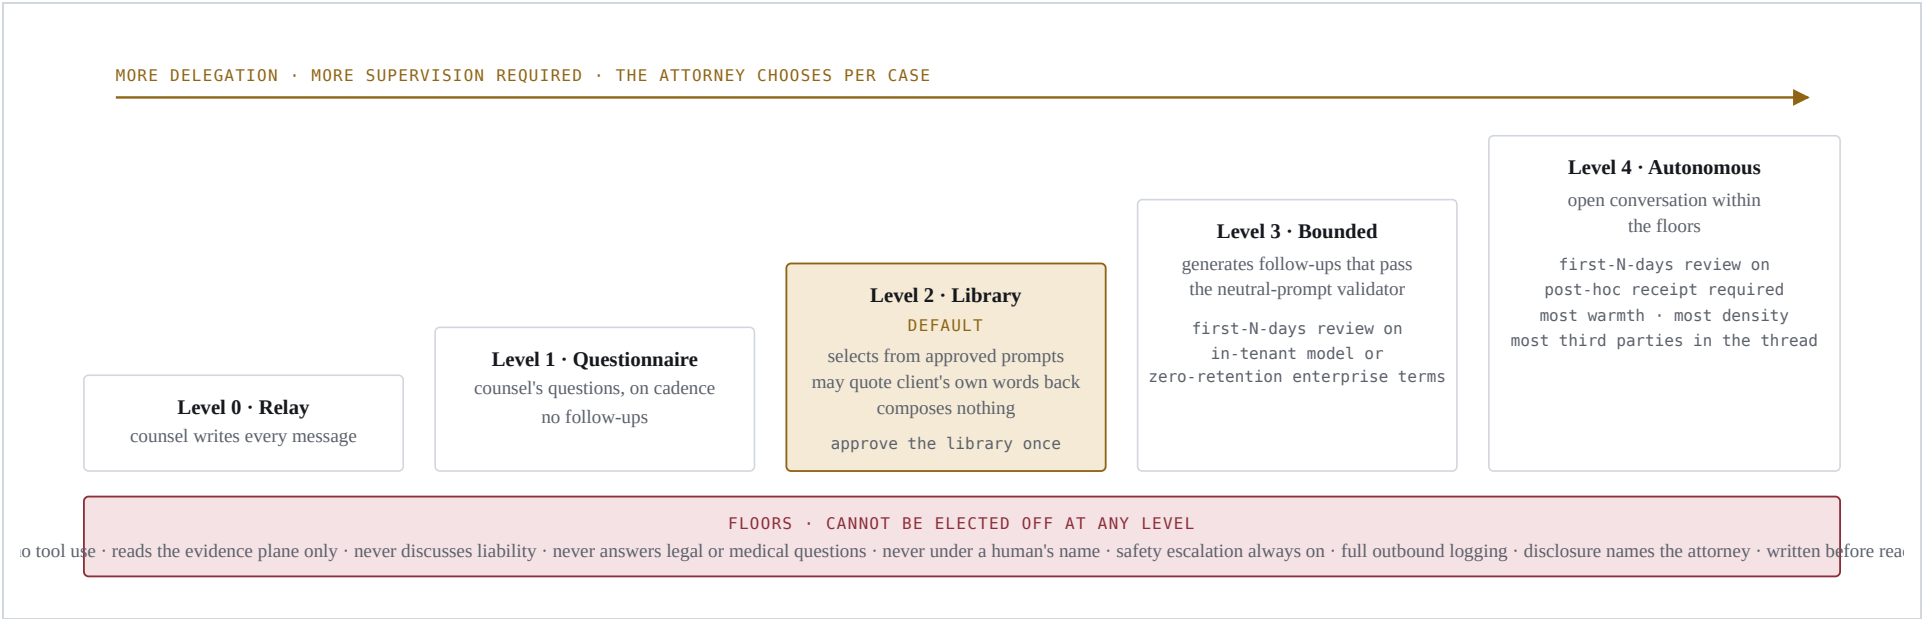
The attorney's week, on a phone



From the thirty-day demo case: four digests opened, seven thread receipts, three attorney dispositions, two attorney messages in the thread, zero approvals pending. About one coffee a week. The Wednesday screen is what "add a thought" became on Mar 25: "Pre-existing-condition point for the damages theory, not a credibility problem on its face: she volunteered it. Discuss with client by phone."

FIGURE 3

How much the attorney delegates, and what never can be



The ladder is the attorney's dial and the supervision instrument at once: choosing a level is a documented act, and the level determines what the digest and approvals screens will ask for. Levels 0–2 keep the fewest third parties in the thread. The floors sit under every level and are not settings.

What "show me how the AI was supervised on this case" produces

One click. Formatted for the two audiences that will ever ask: a disciplinary inquiry and a malpractice carrier. Values below are from the thirty-day synthetic demo case.

SUPERVISION RECORD · MARIA S. (SYNTHETIC)

MAR 2 — APR 1, 2026

Attorney-directed AI assistant · protocol PIRRA-STD-1.4

Supervising attorney	Sara Alvarez
Direction	Authorized Mar 2, 16:41. Level 2 (library with sourced slots). Approval regime: library. Topics authorized and prohibited listed. Forum pinned CA Superior. Letter hash 4c9e... b17a.
Measures	52 of 52 agent messages cite an approved library prompt; 0 composed. 3 quoted-client-words slots validated verbatim. 1 prohibited-topic steer executed (Mar 11). 2 questions routed to a person, 0 answered by the assistant. Assistant had no access to the intelligence plane or the case file.
Client disclosure	Disclosure v4 delivered Mar 3, 09:12 (AI status; attorney named; attorney decides use; monitoring hours; 911; STOP; HUMAN). Consent Mar 3, 09:20.
Knowledge	4 weekly digests, each opened by the supervising attorney (Mar 9, 16, 23, 30). Thread opened 4 times.
Response	9 flags raised, 9 disposed. Attorney: 3 (1 noted, 1 thought added, 1 escalated to the thread with a same-day call). Paralegal: 6. Attorney messages in thread: 2, labeled.
Changes	None. No pause, no mode change, no turn-off in period.
Approvals	None pending at any time.
Exports	None. Impeachment-surface review not yet performed; export blocked.

THE OTHER HALF

What the attorney never has to do

Not asked of the attorney

- **No reply quota.** Receipt is what matters, not how often the attorney writes in the thread.
- **No attestation ritual.** Opening the digest is the receipt.
- **No hunting.** Evidentiary flags come to the phone. Operational noise goes to the paralegal.
- **No policy drafting.** The firm's AI-use policy is generated from settings and kept in sync.
- **No record assembly.** The export above is one click.

Guarded against on the attorney's behalf

- **Handing the transcript to the client before testimony** triggers the §771 warning.
- **Sharing with an expert** triggers a warning.
- **Partial production** is not offered; elections are of the whole record.
- **Exporting before the impeachment-surface review** is blocked, not discouraged.
- **Using a consumer AI tool on the record** is named in the firm policy as the one thing that waives what it touches.

Where each California authority lands

Authority	What it asks	What in the product answers it
Rule 5.3	Measures, authority, knowledge and response over a nonlawyer assistant	Figure 1, entire
Proposed 5.1 (2026)	Firm policies governing the use of AI	Policy generated from firm settings; re-issued on change
Proposed 1.4 Cmt 5	Disclose AI use that materially affects the manner of representation	Disclosure v4 names the AI, the attorney, and who decides use; engagement-letter addendum
Proposed 1.6 "reveal"	Exposure to AI systems that may retain or use the information	Zero-retention terms; in-tenant models for Levels 1–2; customer-managed keys; documented Twilio residuals
May 2026 agentic guidance	Meaningful lawyer review where agents make substantive determinations or transmit externally	Assistant makes no determinations and transmits to no one but the client; flags require human disposition
Rule of Court 10.430	Confidentiality, oversight, disclosure, bias headings for AI policies	Firm policy template on those four headings; compliance snapshot tied to configuration
SB 574 (pending)	No delegating the practice of law to AI; "public generative AI system"	Assistant answers no legal questions and routes them; in-tenant deployment; counsel question on definitions

Companion documents: 01 PIRRA in One Picture and 02 Build Spec. Detail: DR-0005 (autonomy levels and floors), docs/research/california-ethics-governance-2026.md, and the thirty-day demo case in docs/demo-case (timeline.md §3–5 is the supervision record this page renders). Not legal advice; the proposed rules are proposed, and nothing here promises privilege or admissibility.